

Product Expert System — Windows Setup Guide for Beginners

This guide assumes you have a Windows 10 or 11 PC and have never used Docker before.

Every step includes exactly what to click, what to type, and what you should see.

What We're Installing (and Why)

Before we start, here's a plain-English summary of what each piece of software does:

Software	What It Does	Analogy
Docker	Runs our entire application in isolated "containers" — think of it as a	Like running an app inside a
Desktop	virtual computer inside your computer	sealed sandbox
Git	Downloads our project code and tracks changes	Like a smart version of "Download ZIP"
Ollama	Runs AI models locally on your machine — no cloud API needed	Like having a mini ChatGPT on your laptop
VS Code (optional)	A text editor for viewing and editing project files	Like Notepad, but much better for code

Everything else (PostgreSQL database, Redis cache, Nginx web server, Python, FastAPI) **runs inside Docker automatically**. You don't install them separately.

Total cost: \$0. Everything is free for our use.

PHASE 1: Install Required Software

Step 1: Check Your Windows Version

1. Press **Windows key + R** on your keyboard
2. Type `winver` and press **Enter**
3. A window pops up showing your Windows version

You need: Windows 10 version 2004 or higher, OR Windows 11.

If you're on an older version, update Windows first:

Settings → Windows Update → Check for updates.

Step 2: Enable WSL 2 (Windows Subsystem for Linux)

Docker needs WSL 2 to run Linux containers on Windows. This is a one-time setup.

1. Click the **Start** button (Windows icon, bottom-left)
2. Type **PowerShell**
3. Right-click **Windows PowerShell** and choose **Run as Administrator**
4. Click **Yes** when the "Do you want to allow this app..." popup appears
5. In the blue PowerShell window, type this command and press **Enter**:

```
powershell
```

```
wsl --install
```

6. **Wait** — this downloads and installs the Linux subsystem. It may take 5–10 minutes.
7. You'll see messages like:

```
Installing: Virtual Machine Platform
```

```
Installing: Windows Subsystem for Linux
```

```
Installing: Ubuntu
```

8. When it says "**The requested operation is successful**", **restart your computer**.
9. After restart, a terminal window may pop up asking you to create a Linux username and password. You can set these to anything simple (e.g., username: `user`, password: `password`). This is just for the Linux subsystem, not your Windows account.
10. **Verify it worked.** Open PowerShell again (regular, not admin) and type:

```
powershell
```

```
wsl --version
```

You should see something like:

WSL version: 2.x.x

Kernel version: 5.15.x

If you see an error: Go to Settings → Apps → Optional Features → More Windows Features → check both "Virtual Machine Platform" and "Windows Subsystem for Linux" → Restart.

Step 3: Install Docker Desktop

1. Open your web browser and go to: <https://www.docker.com/products/docker-desktop/>
2. Click the big blue "**Download for Windows**" button
3. Open the downloaded file: **Docker Desktop Installer.exe** (Usually in your Downloads folder — about 550 MB)
4. In the installer window:
 - Check "**Use WSL 2 instead of Hyper-V**" (should be checked by default)
 - Check "**Add shortcut to desktop**"
 - Click **OK**
5. **Wait** for the installation to complete (3–5 minutes)
6. Click "**Close and restart**" when prompted
7. After restart, **Docker Desktop** should start automatically.
 - You'll see a whale icon 🐋 in the system tray (bottom-right of your taskbar, near the clock)
 - The whale will be animated while Docker is starting up
 - **Wait until the whale stops moving** — this means Docker is ready
8. If you see a "Docker Subscription Service Agreement" popup:
 - Click **Accept** (it's free for personal and small business use)
9. You may see a "Welcome to Docker Desktop" tutorial — you can **Skip** it.
10. **Verify it worked.** Open PowerShell and type:

```
powershell
```

```
docker --version
```

You should see:

```
Docker version 27.x.x, build xxxxxxxx
```

Then type:

```
powershell  
docker compose version
```

You should see:

```
Docker Compose version v2.x.x
```

If Docker won't start: Make sure WSL 2 is installed (Step 2). Open Docker Desktop → Settings (gear icon) → General → check "Use the WSL 2 based engine" → Apply & Restart.

Step 4: Install Git

1. Open your web browser and go to: <https://git-scm.com/download/win>
2. The download should start automatically. If not, click "**Click here to download manually**"
3. Run the installer: **Git-2.x.x-64-bit.exe**
4. Click through the installer — **use all default options**. Just keep clicking Next → Next → Next → Install
Important defaults to keep:
 - "Git from the command line and also from 3rd-party software"
 - "Use bundled OpenSSH"
 - "Checkout as-is, commit Unix-style line endings" ← **change this one if you see it!** Select "**Checkout as-is, commit as-is**" to avoid line ending issues.
5. Click **Finish**
6. **Close and reopen PowerShell** (so it picks up the new Git installation)
7. **Verify it worked:**

```
powershell  
git --version
```

You should see:

```
git version 2.x.x.windows.x
```

Step 5: Install Ollama (Local AI Engine)

Ollama runs the AI models that power our system's search and Q&A features.

1. Open your web browser and go to: <https://ollama.com/download/windows>
2. Click "Download for Windows"
3. Run the installer: **OllamaSetup.exe**
 - Follow the prompts, click **Install**, then **Finish**
4. Ollama starts automatically as a background service. You'll see a llama icon in your system tray (bottom-right).
5. **Verify it's running.** Open PowerShell and type:

```
powershell  
ollama --version
```

You should see:

```
ollama version 0.x.x
```

6. **Download the AI models we need.** Type each command and wait for it to finish:

```
powershell  
# This downloads the embedding model (~275 MB) — used for search  
ollama pull nomic-embed-text
```

Wait for it to show **success**. Then:

```
powershell  
# This downloads the language model (~4.7 GB) — used for Q&A answers  
ollama pull llama3.1:8b
```

This one takes a while (5–15 minutes depending on your internet speed). Wait for **success**.

7. **Verify both models are installed:**

```
powershell
```

ollama list

You should see both models listed:

NAME	SIZE
nomic-embed-text:latest	274 MB
llama3.1:8b	4.7 GB

8. **Quick test** — make sure the AI is responding:

```
powershell  
curl.exe http://localhost:11434/api/tags
```

You should see a JSON response listing the models. If you see an error, make sure the llama icon is in your system tray.

Step 6: Install VS Code (Optional but Recommended)

VS Code makes it easy to view and edit project files.

1. Go to: <https://code.visualstudio.com/>
2. Click "Download for Windows"
3. Run the installer with default options
4. Check "Add to PATH" during installation

Phase 1 Complete — Checklist

Before continuing, verify everything is installed. Open PowerShell and run each:

```
powershell
```

```
wsl --version      # Should show WSL version 2.x  
docker --version   # Should show Docker version 27.x  
docker compose version # Should show Docker Compose v2.x  
git --version       # Should show git version 2.x  
ollama --version    # Should show ollama version 0.x  
ollama list         # Should show nomic-embed-text and llama3.1:8b
```

If all 6 commands work, you're ready for Phase 2.

PHASE 2: Set Up the Project

Step 7: Create a Project Folder

1. Open **PowerShell** (regular, not admin)
2. Navigate to where you want the project. For example, your Documents folder:

```
powershell  
cd $HOME\Documents
```

3. Create the project folder:

```
powershell  
mkdir product-expert  
cd product-expert
```

4. You're now at: C:\Users\YourName\Documents\product-expert

If your code is in a Git repository, clone it instead:

```
powershell  
cd $HOME\Documents  
git clone https://github.com/your-org/product-expert.git  
cd product-expert
```

Step 8: Place All Project Files

You need to place all the files we've built into this folder. Here's the complete list:

```
product-expert\  
├── .env.example  
├── Dockerfile  
├── docker-compose.yml  
├── requirements.txt  
├── nginx.conf  
├── ops.ps1  
├── Makefile  
├── db-schema.sql  
├── seed.sql  
├── models.py  
├── extraction_pipeline.py  
├── ingestion_orchestrator.py  
├── recommendation_engine.py  
├── rag_retrieval.py  
├── api_layer.py  
├── asynpgsql_repository.py  
├── config.py  
├── alembic.ini  
├── alembic\  
│   ├── env.py  
│   └── versions\  
│       └── 001_initial_schema.py  
└── data\  
    └── samples\  
        └── (your product PDF/TXT/MD files go here)
```

How to create files from the Claude artifacts:

For each artifact I created:

1. Click the **Copy** button on the artifact
2. Open **Notepad** (or VS Code)
3. Paste the content
4. Save with the exact filename shown above
5. **Important:** When saving in Notepad, change "Save as type" to "**All Files (.)**" so it doesn't add **.txt** to the end

Or in VS Code:

1. Open VS Code
2. File → Open Folder → select `(product-expert)`
3. Create each file in the file explorer panel on the left

Critical: Make sure files like `Dockerfile` have NO file extension (not `Dockerfile.txt`). And `docker-compose.yml` ends in `.yml`, not `.yml.txt`.

Step 9: Create the Environment Configuration File

1. Open PowerShell in your project folder:

```
powershell  
cd $HOME\Documents\product-expert
```

2. Copy the template:

```
powershell  
Copy-Item .env.example .env
```

3. Open the `.env` file for editing:

```
powershell  
notepad .env
```

4. The defaults work for local development. The key settings are:

```
env
```

```
# These are already set correctly for Ollama:  
EMBEDDING_PROVIDER=ollama  
EMBEDDING_API_URL=http://host.docker.internal:11434/api/embeddings  
EMBEDDING_MODEL=nomic-embed-text  
LLM_PROVIDER=ollama  
LLM_API_URL=http://host.docker.internal:11434/api/generate  
LLM_MODEL=llama3.1:8b  
  
# Default password (change for production):  
POSTGRES_PASSWORD=expert  
  
# Default API keys for testing:  
API_KEYS=dev-key-001:admin,sales-key-001:sales_engineer,pm-key-001:product_manager
```

5. Save and close Notepad (Ctrl+S, then close)

Step 10: Create the Data Folder for Documents

```
powershell  
  
# Create the folders for storing product documents  
mkdir data  
mkdir data\samples
```

If you have product data sheets (PDFs, text files), copy them into `data\samples\` now. If not, you can add them later.

Step 11: Allow PowerShell Scripts to Run

Windows blocks PowerShell scripts by default. We need to allow our `ops.ps1` script:

```
powershell  
  
Set-ExecutionPolicy -ExecutionPolicy RemoteSigned -Scope CurrentUser
```

Type **Y** and press Enter when asked to confirm.

Phase 2 Complete — Checklist

```
powershell

# Make sure you're in the project folder
cd $HOME\Documents\product-expert

# Check key files exist
Test-Path Dockerfile      # Should say True
Test-Path docker-compose.yml # Should say True
Test-Path .env              # Should say True
Test-Path ops.ps1           # Should say True
Test-Path db-schema.sql     # Should say True
Test-Path seed.sql          # Should say True
Test-Path api_layer.py      # Should say True
```

If all say True, you're ready for Phase 3.

PHASE 3: Build and Start the Application

Step 12: Make Sure Docker Desktop Is Running

1. Look at your system tray (bottom-right of taskbar, near the clock)
2. You should see the Docker whale icon 
3. If it's not there, open Docker Desktop from the Start Menu
4. **Wait until the whale stops animating** (means Docker is ready)

Also verify Ollama is running (llama icon in system tray). If not:

```
powershell

ollama serve
```

Step 13: Build the Docker Containers

This downloads base images and installs all dependencies. **First time takes 5–10 minutes.**

```
powershell
```

```
cd $HOME\Documents\product-expert
```

```
docker compose build
```

What you'll see:

```
[+] Building 3/3
=> [api] Building...
=> => downloading python:3.12-slim
=> => installing pip packages
=> => copying application files
=> [api] DONE
```

If you see errors:

- "Cannot connect to the Docker daemon" → Docker Desktop is not running. Start it and wait for the whale to stop animating.
- "no such file or directory: Dockerfile" → You're not in the right folder. Run `cd $HOME\Documents\product-expert`.
- "failed to solve: python:3.12-slim: error pulling image" → Check your internet connection.

Step 14: Start All Services

```
powershell
```

```
docker compose up -d
```

The `-d` flag means "run in background" (detached mode).

What you'll see:

```
[+] Running 5/5
✓ Network product-expert_backend    Created
✓ Volume "product-expert_pg_data"   Created
✓ Container product-expert-db      Healthy
✓ Container product-expert-redis   Started
✓ Container product-expert-api    Started
✓ Container product-expert-proxy   Started
```

Wait about 30 seconds for everything to fully start up, then check status:

```
powershell
```

```
docker compose ps
```

You should see all 4 containers with Status "Up":

NAME	STATUS	PORTS
product-expert-api	Up (healthy)	0.0.0.0:8000->8000/tcp
product-expert-db	Up (healthy)	0.0.0.0:5432->5432/tcp
product-expert-proxy	Up	0.0.0.0:80->80/tcp
product-expert-redis	Up	0.0.0.0:6379->6379/tcp

If a container shows "Restarting" or "Exit": Check its logs:

```
powershell
```

```
docker compose logs api    # Check the API container  
docker compose logs postgres # Check the database
```

Step 15: Verify Everything Works

Test 1: Health Check

```
powershell
```

```
curl.exe -s http://localhost:8000/health
```

Expected output (something like):

```
json
```

```
{"status":"healthy","components":{"database":"connected","redis":"connected"}}
```

Important: Use `curl.exe` (with .exe), not just `curl`. In PowerShell, `curl` is an alias for something else.

Test 2: Check the Database Was Seeded

```
powershell
```

```
docker compose exec postgres psql -U expert -d product_expert -c "SELECT code, name FROM brands;"
```

Expected output:

```
code | name  
----+  
ABS | American BioTech Supply  
LABRepCo | LABRepCo  
Corepoint | Corepoint Scientific  
Celsius | °celsius Scientific  
CBS | CBS / CryoSafe
```

Test 3: API Responds to Requests

```
powershell
```

```
curl.exe -s http://localhost:8000/use-cases -H "X-API-Key: dev-key-001"
```

Expected: A JSON list of use cases (vaccine_storage, laboratory_general, etc.)

Test 4: Open in Your Browser

Open your web browser and go to: <http://localhost:8000/docs>

You should see the **FastAPI Swagger UI** — an interactive API documentation page where you can test all endpoints.

🎉 THE SYSTEM IS RUNNING!

You should now have:

- **API** at <http://localhost:8000>
 - **API docs** at <http://localhost:8000/docs>
 - **Frontend** at <http://localhost> (via Nginx — needs frontend files in `(frontend\dist\)`)
-

PHASE 4: Ingest Documents and Test

Step 16: Ingest a Product Document

Place a product PDF in `(data\samples\)`, then:

```
powershell
```

```
# Ingest a single file  
\ops.ps1 ingest-file data\samples\ABS_Premier_26S_PDS.pdf  
  
# Or ingest all files in the data folder  
\ops.ps1 ingest-sample
```

Expected output:

```
json  
  
{  
  "job_id": "a1b2c3d4...",  
  "status": "processing",  
  "files_accepted": 1,  
  "message": "1 file(s) queued for processing"  
}
```

Step 17: Run the Full Smoke Test

```
powershell  
  
\ops.ps1 smoke-test
```

This runs 6 checks in sequence: health, stats, use-cases, product search, conflicts, and recommendation. Each should return JSON data without errors.

Step 18: Try the Q&A Feature

```
powershell  
  
\ops.ps1 test-ask "What vaccine refrigerators meet NSF 456?"
```

Step 19: Check for Spec Conflicts

```
powershell  
  
\ops.ps1 conflicts
```

If there are conflicts (from overlapping data in documents), resolve them:

```
powershell
```

```
# Via the API directly  
curl.exe -X PUT http://localhost:8000/conflicts/CONFLICT-ID-HERE -H "X-API-Key: dev-key-001" -H "Content-Type: app
```

PHASE 5: Daily Operations

Starting the System (After a Reboot)

When you restart your computer:

1. **Docker Desktop** should start automatically (check for the whale icon)
 - If not, open it from the Start Menu
2. **Ollama** should start automatically (check for the llama icon)
 - If not, open PowerShell and run `ollama serve`
3. Start the application:

```
powershell  
cd $HOME\Documents\product-expert  
.ops.ps1 up
```

Stopping the System

```
powershell  
.ops.ps1 down
```

Viewing Logs

```
powershell  
.ops.ps1 logs      # API logs only  
.ops.ps1 logs-all  # All services
```

Press **Ctrl+C** to stop watching logs.

Common ops.ps1 Commands Quick Reference

```
powershell
```

```
\ops.ps1 help      # Show all commands
\ops.ps1 up        # Start everything
\ops.ps1 down      # Stop everything
\ops.ps1 status     # Check what's running
\ops.ps1 health    # Quick health check
\ops.ps1 stats      # System statistics
\ops.ps1 logs       # View API logs
\ops.ps1 smoke-test # Run all tests
\ops.ps1 ingest-file <path> # Ingest a document
\ops.ps1 products    # Search all products
\ops.ps1 conflicts   # Check spec conflicts
\ops.ps1 db-shell     # Open database terminal
\ops.ps1 db-dump      # Backup database
```

Troubleshooting

"docker: command not found" or "not recognized"

- Docker Desktop is not installed or not in your PATH.
- Restart PowerShell after installing Docker Desktop.
- Make sure Docker Desktop is running (whale icon in system tray).

"Cannot connect to the Docker daemon"

- Docker Desktop is installed but not running.
- Click the Docker Desktop icon to start it.
- Wait for the whale to stop animating before running commands.

"curl.exe: command not found"

- You're on an older Windows version. Install curl:

```
powershell  
winget install cURL.cURL
```

"port is already in use" or "address already in use"

Another program is using port 8000, 5432, 6379, or 80.

```
powershell
```

```
# Find what's using port 8000  
netstat -ano | findstr :8000  
  
# The last column is the PID. Kill it:  
taskkill /PID 12345 /F
```

Common culprits: local PostgreSQL (port 5432), IIS/Apache (port 80), or another Docker project.

Container keeps restarting

```
powershell  
  
# Check what's wrong  
docker compose logs api --tail 50  
  
# Common fixes:  
# 1. Database not ready yet — wait 30 seconds and check again  
# 2. Wrong .env settings — verify POSTGRES_PASSWORD matches  
# 3. Port conflict — see above
```

"Ollama connection refused" in API logs

The API container can't reach Ollama on your host machine.

```
powershell  
  
# 1. Make sure Ollama is running  
curl.exe http://localhost:11434/api/tags  
  
# 2. Make sure .env has the right URL  
# It should be: http://host.docker.internal:11434/api/embeddings  
# NOT: http://localhost:11434/api/embeddings
```

"no matching manifest for windows/amd64"

You need to switch Docker to Linux containers: → Right-click the Docker whale in system tray → **"Switch to Linux containers..."**

"File not found" errors during build

Make sure your files don't have hidden `.txt` extensions. In File Explorer: → View tab → check **"File name extensions"** to see real extensions.

Everything was working but now it's broken

```
powershell
```

```
# Nuclear option: full reset (deletes all data!)
```

```
.\ops.ps1 reset
```

What's Running Where (Summary)

After `\ops.ps1 up`, you have 4 containers running:

